


Effectiveness Evaluation

Manual versus agent-assisted comparison

Purpose. Compare the two workflows using the same scenarios and separate correctness and control from efficiency gains.

Paired evaluation design



Use paired scenarios so any improvement can be attributed to the workflow rather than easier cases.

1. Paired comparison design

Each selected scenario is performed twice: once using the manual baseline workflow and once using the agent-assisted workflow. Inputs and expected answers are held constant. If multiple testers are available, alternate the order of manual and agent runs to reduce learning effects.

2. Primary evaluation table

Metric	Manual	Agent	Difference	Comment
Preparation time (s)				
Review time (s)				
Total task time (s)				
Manual actions				

Metric	Manual	Agent	Difference	Comment
Raw-file reopens				
Correct decision				
Correct quantity				
Blocking exceptions found				
PO field completeness				
Traceability complete				

3. Calculation rules

Measure	Formula / rule
Time reduction %	$(\text{Manual time} - \text{Agent time}) / \text{Manual time} \times 100$
Action reduction %	$(\text{Manual actions} - \text{Agent actions}) / \text{Manual actions} \times 100$
Error reduction	Manual error count - Agent error count
Review reduction %	$(\text{Manual review time} - \text{Agent review time}) / \text{Manual review time} \times 100$
Exception recall	Caught expected blockers / expected blockers
Net benefit gate	Efficiency claims are reported only when correctness / control criteria pass.

4. “Gate, then optimise” interpretation

The evaluation should not average speed and correctness into one score. First check the safety / correctness gates. Only if the agent meets them should time and action savings be treated as a benefit.

Stage	Requirement
Gate 1 - correctness	Expected decision and quantity outcome matches the agreed answer.
Gate 2 - exception integrity	Blocking cases remain visible and no invalid line is silently promoted.
Gate 3 - approval integrity	Human approval and post-edit review behaviour are preserved.
Gate 4 - traceability	Reviewer can identify why a line exists and where it came from.
Optimise after gates	Compare time, actions, review effort and rework.

5. Result narrative template


Example structure for the final report / demo after measurements are available:

“Across [N] paired scenarios, the agent-assisted workflow reduced median preparation time from [A] s to [B] s and reduced median manual actions from [C] to [D]. Decision and quantity accuracy remained at [X%], while [Y/Y] expected blocking exceptions were surfaced. Reviewers reopened raw source files [M] times manually versus [N] times with the agent. The largest remaining friction was [issue], mainly in [scenario type].”

6. What not to claim

- Do not claim cost savings without an agreed cost model.
- Do not claim production-grade reliability from a small synthetic demo set.
- Do not claim the agent “eliminates human error”; it changes where human judgement is needed.
- Do not claim supplier optimisation if the prototype uses a preset default supplier.
- Do not claim ERP integration if the output is only a prototype draft / export.

Suggested demo evidence sequence



The demo should prove both automation and control. A fast system that hides errors is not a successful result.